



## **PROPOSAL FOR AGENT-0 CHAIRMAN AI TWIN PROJECT**

**Prepared by:** Dynamiq Technologies, Inc.

**Prepared for:** Abu Dhabi Department of Government Enablement

**Date:** December 17, 2025

<b>1. Executive Summary</b>	<b>4</b>
1.1 Strategic Value	4
1.2 Expected Outcomes & Success Metrics	5
1.3 Why Dynamiq	6
<b>2. About Dynamiq</b>	<b>7</b>
<b>3. The Agent-0 Vision</b>	<b>9</b>
3.1 A Framework for Leadership	9
3.2 The Three-Dimensional Persona Model	9
<b>4. Sovereign Deployment Overview</b>	<b>10</b>
<b>5. Proposed Solution Architecture (UAE-Centric Hybrid Cloud)</b>	<b>11</b>
5.1 Architectural Principles	11
5.2 On-Premises GPU Infrastructure for Sovereign AI	11
5.3 Proposed Hardware: 2x Servers with 8x NVIDIA H200 GPUs	11
5.4 AI Component Requirements	12
5.5 Recommended Deployment Configuration	12
<b>6. The Agent-0 Decision Engine</b>	<b>13</b>
6.1 The 3-Level Authority Framework	13
6.2 Decision Process Flow	15
6.3 The Authority Barometer	16
<b>7. The Institutional Knowledge Graph</b>	<b>17</b>
7.1 Knowledge Graph Architecture	17
7.2 From Data to Wisdom	17
<b>8. Project Team Structure</b>	<b>18</b>
8.1 Key Personnel and Responsibilities	18
<b>9. Project Timeline, Workstreams &amp; Deliverables</b>	<b>19</b>
9.2 Workstream Activities by Phase	19
<b>10. Assumptions and Dependencies</b>	<b>20</b>
10.1 Key Assumptions	20
10.2 Client Dependencies	20
<b>11. Commercial Proposal</b>	<b>21</b>
11.1 Platform Licensing Fee	21
11.2 Implementation Cost	22
<b>Appendix A – Dynamiq Technical Capabilities</b>	<b>24</b>
12.1 Sovereign AI Infrastructure & Deployment Flexibility	24
12.2 Low-Code Workflow & Agent Builder	25
12.3 Advanced RAG and Knowledge Management	26
12.4 Complex Multi-Agent Orchestration and Tool Automation	27
12.5 Native, Governed Tool Calling with Full Auditability	28
12.6 Human-in-the-Loop	29
12.7 Comprehensive Agent Evaluation & Observability	30
12.8 Execution Guardrails and Enterprise Policy Enforcement	31

12.9 Open-Source LLM Deployment and Fine-Tuning	32
12.10 API-First Integration Architecture	34
12.11 Pre-Built UI Components and LLM Gateway	35
12.12 Open Architecture and Extensibility	37
<b>Appendix B – Dynamiq Case Studies &amp; Tutorials</b>	<b>39</b>
<b>Appendix C – Dynamiq Infrastructure Components Info</b>	<b>40</b>

# 1. Executive Summary

Dynamiq is honored to present this proposal for the creation of **Agent-0**, an AI-powered platform designed to reimagine executive leadership within the Abu Dhabi Department of Government Enablement (DGE). This document outlines a comprehensive, enterprise-grade solution to be deployed on a secure and sovereign **Amazon Web Services (AWS)** cloud infrastructure. It details the complete project scope, from technical architecture and implementation to the project team and a transparent commercial framework.

Agent-0 is designed to address three core constraints of modern leadership: **bandwidth, speed, and fragmented context**. In an environment where critical information is scattered across systems and stakeholders, and decision cycles keep accelerating, executives are forced to make high-impact calls without consistent access to the full picture. By creating a sophisticated AI-powered digital twin, the platform will amplify the reach, consistency, and effectiveness of leadership. It will enable high-quality decision-making and stakeholder engagement at a scale previously unattainable. This initiative is not about simple task automation; it represents a fundamental augmentation of leadership itself.

The platform is built on five synergistic systems:

1. **Dynamic Decision Architecture:** To intelligently analyze and manage the lifecycle of executive decisions through a three-tier authority framework with a risk-aware Authority Barometer.
2. **Comprehensive Context Engine:** To synthesize real-time, organization-wide knowledge and external intelligence into an actionable, always-on intelligence layer with emotional intelligence capabilities.
3. **Multi-Channel Engagement Layer:** To manage authentic and personalized stakeholder interactions at scale across all engagement tiers, including organizational-wide employee engagement.
4. **Digital Persona Framework:** To learn and authentically replicate the executive's unique leadership style, values, and voice with multi-modal, simultaneous presence capability.
5. **Reinforcement Learning System:** To ensure the platform continuously learns and improves from every interaction, outcome, and human override, with explicit ethical guardrails.

Dynamiq's experience in building sovereign, agentic AI platforms, combined with our strategic partnership with AWS, makes us the ideal partner to bring the vision of Agent-0 to life. We are confident this initiative will deliver transformative value to the DGE and establish Abu Dhabi as a global leader in AI-driven governance.

## 1.1 Strategic Value

Agent-0 delivers significant value by addressing the most critical challenges faced by modern executives:

<b>Challenge</b>	<b>Agent-0 Platform</b>	<b>Strategic Value</b>
<b>Decision Bottlenecks</b>	Three-tier decision authority framework with AI-powered analysis, automated workflows, and prioritized briefings.	Accelerates decision-making, improves the quality and consistency of choices, and reduces organizational risk.
<b>Fragmented Engagement</b>	A scalable digital persona capable of simultaneous, multi-modal interactions across all stakeholder tiers.	Deepens stakeholder relationships and ensures strategic alignment.
<b>Information Overload</b>	Real-time synthesis of all organizational data and external intelligence into a single, always-on context engine with emotional intelligence.	Provides a definitive source of truth for data-driven leadership and surfaces critical insights.
<b>Inconsistent Execution</b>	A consistent framework for applying the executive's vision and decision-making logic with continuous learning and ethical guardrails.	Ensures strategic alignment and consistent execution of leadership's intent across the organization.

## 1.2 Expected Outcomes & Success Metrics

The success of the Agent-0 platform will be measured against the specific Key Performance Indicators (KPIs) outlined in the RFP. Our implementation is designed to meet or exceed these ambitious targets.

<b>KPI Category</b>	<b>RFP Target</b>	<b>How Agent-0 will achieve it</b>
<b>Decision Throughput</b>	10-50x increase	By automating Level 1 decisions and streamlining Level 2 approvals, the platform will dramatically increase the number of decisions that can be processed.
<b>Decision Cycle Time</b>	70% reduction	The platform will reduce the time it takes to gather information, get approvals, and communicate decisions, leading to a significant reduction in cycle time.
<b>Decision Quality</b>	25% improvement	By providing comprehensive data, flagging risks, and learning from past outcomes, the platform will improve the quality and consistency of decision-making.
<b>Stakeholder Responsiveness</b>	95% of queries	The Multi-Channel Engagement Layer and Digital Avatar will be able to provide instant, accurate responses to a wide range of stakeholder inquiries.

KPI Category	RFP Target	How Agent-0 will achieve it
	answered within 2 hours	
<b>Auditability</b>	100% Auditability	The immutable audit trail will provide a complete, verifiable record of every decision and action taken by the platform.

### 1.3 Why Dynamiq

Dynamiq is the right partner for this project for several key reasons:

- **Proven platform:** We have a mature, enterprise-grade agentic OS for building complex multi-agent systems.
- **Deep AWS expertise:** As an Advanced AWS Partner, we have certified expertise in deploying secure and high-performance AI solutions on the AWS cloud.
- **Sovereign AI experience:** We have a proven track record of delivering sovereign AI solutions for government clients, ensuring the highest standards of security and data residency.
- **Multi-disciplinary team:** We bring together a team of AI engineers, data scientists, cloud architects, and change management experts to ensure a successful implementation.

## 2. About Dynamiq

Dynamiq is a sovereign-first agentic operating system built for regulated environments. Its low-code visual builder and modular architecture enable enterprises to design, test, and deploy single-agent or multi-agent systems on-premises, in the cloud, or in hybrid/air-gapped environments. Dynamiq unifies agent orchestration, retrieval-augmented generation (RAG), and workflow management in a single stack - so teams can build complex multi-agent workflows, connect domain data through enterprise connectors and RAG pipelines, and monitor performance with built-in observability and evaluations. Dynamiq also provides enterprise-grade guardrails, role-based access control (RBAC), and no-code human-in-the-loop approval steps, ensuring agents remain auditable and compliant while still benefiting from autonomous planning and execution. The platform is underpinned by Dynamiq's open-source, code-first agentic framework (Apache 2.0), enabling transparent customization and interoperability.

Dynamiq delivers significant economic and operational benefits by consolidating capabilities that typically require 15+ separate tools, compressing development cycles from ~6 months to ~14 days while reducing build. Enterprises can rapidly prototype and scale agent workflows, and use 1,000+ pre-built integrations to automate document ingestion, tool calling, and API-driven processes.

Security and sovereignty are core to Dynamiq's design. The platform is deployed in the customer's own infrastructure, with dedicated environments, fine-grained access controls, audit trails, and monitoring. Dynamiq includes policy-based guardrails, PII leakage detection, and prompt-injection filtering to help ensure outputs meet enterprise compliance requirements (SOC 2, HIPAA, GDPR), while supporting strong data-residency and governance needs for public-sector and sensitive workloads.

Dynamiq is backed by a seasoned team with deep MLOps and enterprise AI experience. Founder and CEO Vitalii Duk previously led MLOps platforms at OLX and Careem (Uber) and served as CTO at ADQ.

### Key Highlights

- Dynamiq's low-code and no-code platform is built on top of our [open-source, Apache-2.0 licensed agentic framework](#), enabling engineers to build, extend, and deploy production AI agents quickly in a code-first manner.
- We are a [global technology partner of IBM in the agentic AI space](#), and an early contributor to IBM's agent ecosystem initiatives, including Agent Connect and the Watsonx Orchestrate Agent Catalog.
- Dynamiq is an AWS and Microsoft Azure partner, with listings across [AWS Marketplace](#) and [Azure Marketplace](#) to support streamlined procurement and enterprise deployments.
- We are a contributor to [Cisco's AGNTCY initiative](#), now under [Linux Foundation governance](#), helping [define open standards for interoperable, secure multi-agent systems](#).

- Dynamiq is a [Cloudera partner](#), enabling enterprises to pair agentic workflows with a robust hybrid data ecosystem across cloud and on-prem environments.
- Dynamiq's platform is built for regulated environments and maintains a compliance-ready foundation aligned with SOC 2, HIPAA, and GDPR requirements.

### 3. The Agent-0 Vision

#### 3.1 A Framework for Leadership

We have aligned the Agent-0 vision with the CEO Excellence framework. The platform is designed to augment the executive's capabilities across all six critical dimensions of leadership.

Leadership Area	How Agent-0 augments
<b>1. Executive Engagement</b>	The Meeting and Personal Assistants will streamline meeting prep, automate follow-ups, and manage routine approvals.
<b>2. Corporate Strategy</b>	The platform will provide deep analytical support for strategic planning, including scenario modeling and risk assessment.
<b>3. External Stakeholders</b>	Agent-0 will monitor external communications and news to help manage public perception and stakeholder relationships.
<b>4. Organization Alignment</b>	By consistently applying the executive's decision-making framework, the platform will ensure strategic alignment.
<b>5. Teams and Processes</b>	The platform will help optimize team performance by identifying bottlenecks and suggesting improvements.
<b>6. Personal Working Norms</b>	Agent-0 will learn and adapt to the executive's personal working style for a seamless experience.

#### 3.2 The Three-Dimensional Persona Model

To create an authentic digital twin, we will use a three-dimensional persona model to capture the executive's observable behaviors, cognitive processes, and core values.

Dimension	Description	Data Sources
<b>The Sensory Persona (Observable Self)</b>	The executive's appearance, voice, and communication style.	3D scans, video footage, audio recordings, and written communications.
<b>The Cognitive Engine (Executing Self)</b>	How the executive thinks, acts, and makes decisions.	Decision logs, meeting transcripts, and calendar data.
<b>The Core Identity (Guiding Self)</b>	The executive's core values, mission, and vision.	Public statements, vision documents, and structured interviews.

## 4. Sovereign Deployment Overview

Agent-0 is deployed **hybrid-first with a sovereign center of gravity in DGE-controlled infrastructure**. The on-prem (and optionally air-gapped) environment hosts Dynamiq's agentic OS, including the **multi-agent orchestration and decision engine**, sovereign knowledge components (e.g., vector/graph stores), and **locally served open-source LLMs** for sensitive and confidential workloads. Dynamiq enforces **RBAC**, encrypts data **at rest and in transit**, and maintains **end-to-end audit logs** for every agent step, tool call, and decision.

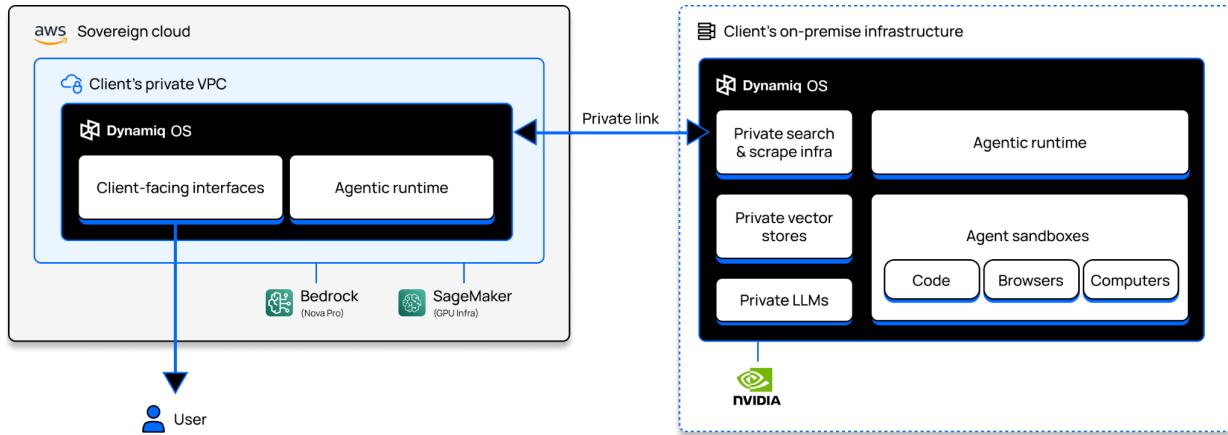


Figure 4.1. Hybrid sovereign architecture: Dynamiq runs in the client's on-prem environment with optional AWS UAE VPC services connected via private link.

External connectivity is **explicitly scoped and policy-controlled**. When Agent-0 needs approved, non-sensitive enrichment (e.g., public sources, whitelisted third-party services, or UAE-hosted cloud capabilities), the orchestrator routes requests to **approved endpoints in the AWS UAE region (me-central-1)** using private connectivity patterns (e.g., PrivateLink / VPN / Direct Connect, as applicable). **Only the minimum required data** leaves the sovereign boundary (with redaction and policy filtering where configured), and results are returned to the on-prem environment so that **authoritative processing, storage, and governance remain sovereign**.

Operations teams maintain full visibility through built-in observability that traces **model selection, retrieval steps, tool invocations, latencies, and outcomes** across on-prem and cloud hops. This architecture keeps the on-prem deployment as the **system of record**, while enabling selective use of UAE cloud services for **scalability and controlled external integrations**, without compromising confidentiality or auditability.

## 5. Proposed Solution Architecture (UAE-Centric Hybrid Cloud)

Our architecture is a secure, resilient, and scalable hybrid model that strictly adheres to the client's requirements for data sovereignty and the use of AWS services available in the UAE (me-central-1) region.

### 5.1 Architectural Principles

The architecture is guided by four key principles: Hybrid by Design, ensuring sensitive data and core processing remain on DGE's on-premises infrastructure while leveraging AWS for scalability; UAE-First, with all AWS services verified to be available in the me-central-1 region; Security-in-Depth, implementing a multi-layered security approach built on a Zero-Trust model; and Powered by Dynamiq's Agentic OS, with our proprietary operating system providing the core scaffolding for agent orchestration, data management, and model deployment.

### 5.2 On-Premises GPU Infrastructure for Sovereign AI

To ensure full data sovereignty and provide maximum performance for the AI models at the heart of Agent-0, we propose a powerful on-premises GPU infrastructure to be housed within DGE's private data center. This infrastructure will be dedicated to the most sensitive AI workloads, including the fine-tuning of the Chairman's Persona and the real-time inference of the primary LLM.

### 5.3 Proposed Hardware: 2x Servers with 8x NVIDIA H200 GPUs

We recommend a two-server configuration, each equipped with 8x NVIDIA H200 GPUs. This provides a formidable total of 16 H200 GPUs and 2,256 GB of high-bandwidth VRAM, offering more than enough capacity for the entire Agent-0 workload with significant room for future expansion.

<b>Server</b>	<b>GPUs</b>	<b>VRAM per GPU</b>	<b>Total VRAM</b>
Server 1	8x H200	141 GB	1,128 GB
Server 2	8x H200	141 GB	1,128 GB
<b>Total</b>	<b>16x H200</b>		<b>2,256 GB</b>

## 5.4 AI Component Requirements

Component	VRAM Required	GPU Requirement
Primary LLM	~670 GB	8 GPUs minimum
Text-to-Speech	4-6 GB	1 GPU
Speech-to-Text	4-6 GB	1 GPU
Avatar Generation	45-80 GB	1 GPU

## 5.5 Recommended Deployment Configuration

All 8 H200 GPUs will be dedicated to running the primary instance of the foundation LLM using 8-way tensor parallelism. This configuration maximizes the performance of the core LLM, enabling the highest throughput and lowest latency for real-time decision support.

This recommended setup ensures that the primary LLM has dedicated resources for maximum performance, while still maintaining full redundancy and dedicated resources for the critical voice and avatar components. The entire on-premises deployment will be managed by Dynamiq's Agentic OS on-premise runtime, ensuring seamless integration with the AWS components of the hybrid cloud architecture.

For the recommended configuration, inter-server networking requires a minimum of 2x 100GbE between servers for service discovery, load balancing, failover coordination, and shared storage access.

## 6. The Agent-0 Decision Engine

The heart of Agent-0 is its sophisticated Decision Engine, which moves beyond simple automation to provide nuanced, context-aware decision support. It is designed to handle the complexity of executive decision-making through a structured, multi-level framework.

### 6.1 The 3-Level Authority Framework

The engine classifies every decision request into one of three authority levels, ensuring that the right degree of human oversight is applied.

Level	Type	Description	Example
Level 1	Autonomous Agent	For high-frequency, low-risk operational decisions governed by clear rules. The agent can execute these decisions autonomously with a full audit log.	Routine procurement approvals within pre-defined budget limits; scheduling standard meetings.
Level 2	Collaborative Agent	For strategic decisions with moderate impact. The agent analyzes the situation, generates recommendations and scenarios, but requires human approval to execute.	Resource allocation for a new project; responding to a partnership proposal.
Level 3	Advisory Agent	For high-stakes, complex decisions requiring deep executive judgment. The agent acts as a strategic advisor, providing a comprehensive briefing, data analysis, and risk assessment, but the final decision and execution are fully human-led.	Major investment decisions; organizational restructuring.

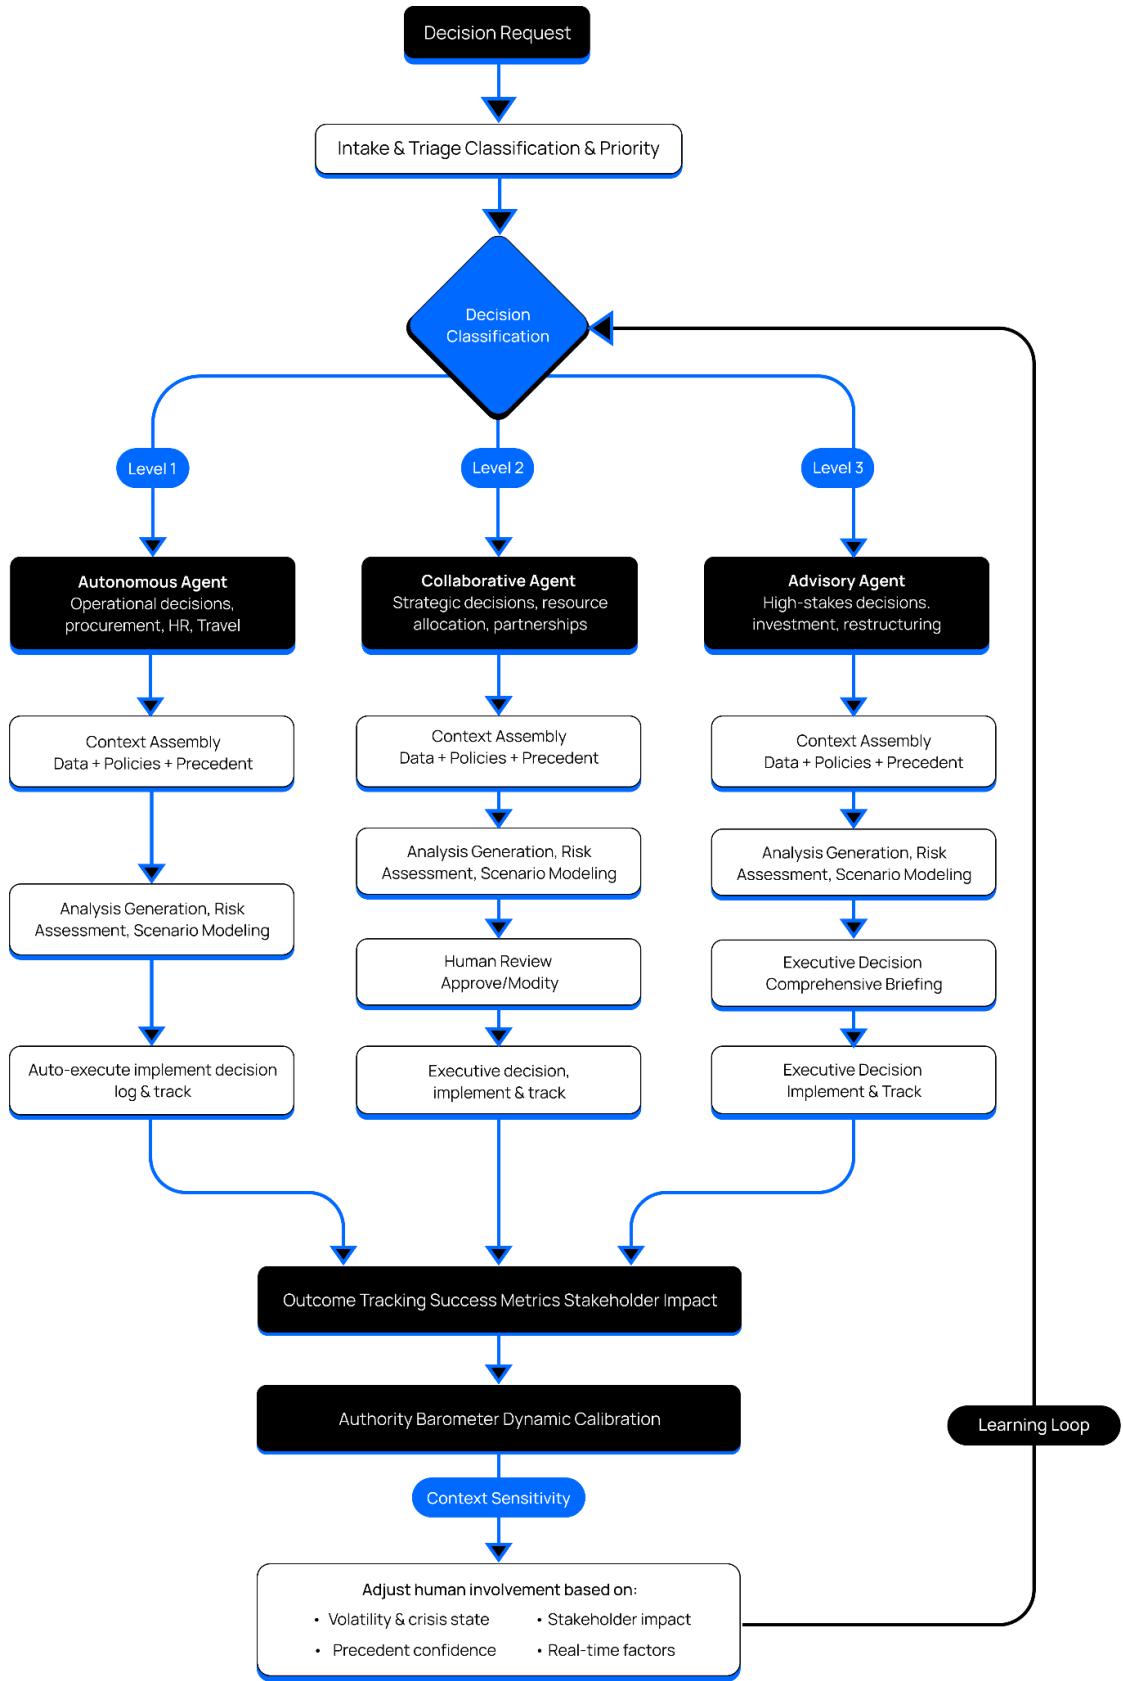


Figure 6.1.1. 3-level authority framework and Authority Barometer

## 6.2 Decision Process Flow

The decision process follows five key stages:

1. **Intake & Triage** ingests decision requests from any channel (email, meeting, ECAS, etc.).
2. **Classification & Priority Agent** is used to determine the decision type and urgency, assigning it to the appropriate authority level.
3. **Context Assembly Agent** gathers all relevant data - historical precedents, policies, stakeholder information, and real-time data - from the Institutional Knowledge Graph.
4. **Analysis & Generation Agent** performs a risk assessment, models potential scenarios, and generates a recommended course of action or a comprehensive briefing document.
5. **Execution & Tracking Agent** either executes the decision autonomously, presents it for approval, or provides advisory support, with all outcomes tracked.

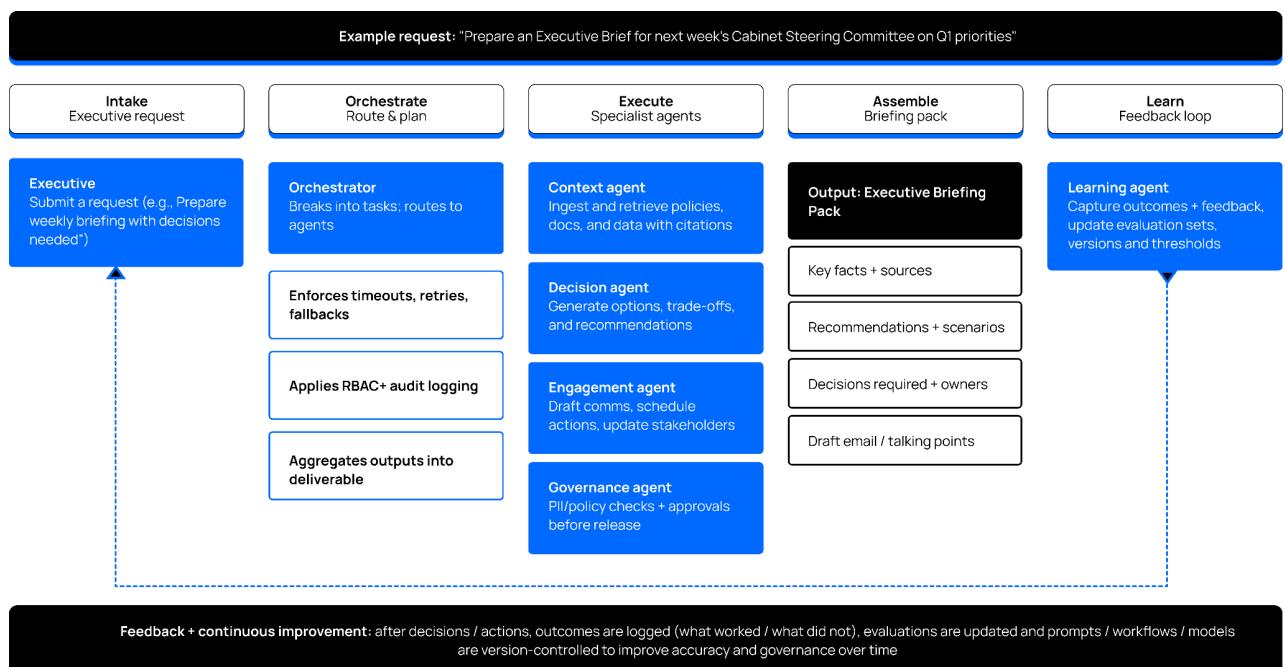


Figure 6.2.1. 5-stage multi-agent workflow

### 6.3 The Authority Barometer

A key innovation is the **Authority Barometer Agent**, a dynamic, agentic mechanism designed to adjust the level of autonomy in real-time.

This sophisticated agent continuously monitors a wide range of high-risk signals, including market volatility, indicators of a crisis state, potential stakeholder impact, and the confidence level of existing operational precedents.

If the agent detects a significant increase in risk or uncertainty, it possesses the capability to automatically escalate a decision (for example, shifting it from a low-autonomy Level 1 to a higher-oversight Level 2). This crucial function ensures that critical human judgment and executive review are injected precisely when the situation demands it most.

This intelligent, agent-based control system is essential for maintaining optimal operational speed and efficiency while simultaneously safeguarding the organization against unforeseen risks and potential negative consequences.

## 7. The Institutional Knowledge Graph

The platform's long-term memory and contextual understanding are powered by the Institutional Knowledge Graph, a dynamic network of interconnected data managed by Dynamiq's Agentic OS and built on Amazon Neptune and local graph databases hosted in a sovereign environment.

### 7.1 Knowledge Graph Architecture

The graph connects the organization's most critical assets—people, decisions, policies, and outcomes - into a coherent, queryable whole.

Core Entities (Nodes) include person, organization, decision, policy, project, meeting, document, and risk. These entities are connected by rich relationships (edges) such as ATTENDED, MADE\_DECISION, BASED\_ON, IMPACTS, and OWNS\_RISK, creating a rich tapestry of institutional knowledge.

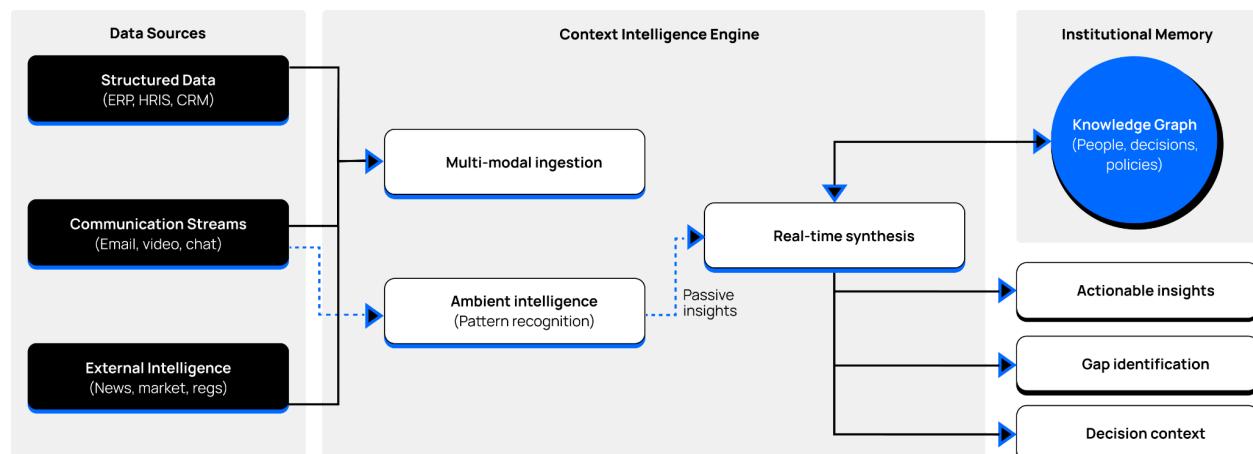
The graph is dynamic and self-updating, with the context agent continuously ingesting new information from integrated systems, automatically creating new entities and relationships.

### 7.2 From Data to Wisdom

The Knowledge Graph transforms raw data into actionable wisdom by enabling complex queries that are impossible with traditional databases.

Examples include Precedent Analysis ("show me all past decisions related to infrastructure projects in Al Ain, who was involved, and what the outcomes were"), Stakeholder Impact Analysis ("if we change this HR policy, which teams and projects will be most affected based on past dependencies?"), and Expert Finding ("who has the most experience and best track record in negotiating international trade agreements?").

This capability ensures that every new decision is informed by the full weight of the organization's cumulative experience.



## 8. Project Team Structure

### 8.1 Key Personnel and Responsibilities

#	Role	Responsibilities
1	Project Director, Strategy	Executive oversight, stakeholder management, strategic decisions, risk management
2	Principal Technical Architect	System architecture, technology decisions, integration design, technical leadership
3	Lead AI Engineer	AI strategy, LLM implementation, persona development, model optimization
4	Project Manager	Day-to-day delivery, sprint planning, resource management, client communication
5	Business Analyst	Requirements management, UAT coordination, documentation, training
6	Senior AI Engineer	LLM fine-tuning, prompt engineering, AI pipeline development
7	Senior Data Engineer	Data architecture, ETL pipelines, integration development
8	Security Engineer	Security architecture, compliance, audit preparation
9	DevOps Engineer	CI/CD, infrastructure automation, deployment
10	UI/UX Designer	Responsible for the design and implementation of user interfaces.
11	Software Engineer	Engaged in both back-end and front-end software development.
12	QA Engineer	Test planning, automation, performance testing

## 9. Project Timeline, Workstreams & Deliverables

### 9.2 Workstream Activities by Phase

<b>Phase 1: Foundation &amp; POC (Months 1-3)</b>		
<b>Workstream ID</b>	<b>Workstream Name</b>	<b>Key Deliverable in Phase 1</b>
WS1	Infrastructure & Platform	On-prem GPU cluster setup, AWS UAE environment provisioning, Dynamiq agentic OS deployment, security framework implementation.
WS2	Decision Engine	Decision Classification Agent v0.1: A working agent that can categorize incoming requests and apply basic routing rules from the DOA matrix.
WS3	Context Engine	Knowledge Graph Foundation: Initial deployment of the graph and vector databases. Semantic Layer v0.1 with a functional Oracle Fusion Connector POC.
WS4	Engagement & UI	Basic UI/Dashboard: A functional web interface showing the decision queue and agent activity.
WS5	Governance & Testing	POC Demo & Feedback Cycle: A live demonstration of an end-to-end decision workflow, from intake to classification, to gather early stakeholder feedback.

<b>Phase 2: Full Build &amp; Production (Months 4-6)</b>		
<b>Workstream ID</b>	<b>Workstream Name</b>	<b>Key Deliverables in Phase 2</b>
WS1	Infrastructure & Platform	Production hardening, performance optimization, and high-availability configuration.
WS2	Decision Engine	Advanced Analysis Agents: Agents capable of scenario modeling and risk assessment. Full Authority Routing with dynamic adjustments.
WS3	Context Engine	Full Data Ingestion Pipeline: Integration with all remaining data sources (ECAS, Email, Calendar, Documents). Full Vector Search capabilities.
WS4	Engagement & UI	Chairman Persona Development: Communication Style Agents trained and deployed. Voice Agent Capabilities enabled. Full UI/Dashboard with advanced analytics.
WS5	Governance & Testing	Full UAT Cycle, security audit, performance testing, user training, and operational handover.

## 10. Assumptions and Dependencies

### 10.1 Key Assumptions

Category	Assumption
Access	DGE will provide timely access to Oracle Fusion, ECAS, and other systems within 2 weeks of project start
Data	Historical data (emails, decisions, documents) is available and can be shared for AI training
Stakeholders	Chairman will be available for weekly 1-hour sessions for persona development and feedback
Infrastructure	Cloud infrastructure (AWS/Azure UAE region) will be provisioned within 2 weeks of project start
Security	Security requirements and compliance standards will be finalized during Phase 1
Resources	DGE will assign a dedicated Product Owner and technical liaison for the project duration
Decisions	Key decisions (architecture, security, scope) will be made within 5 business days of presentation

### 10.2 Client Dependencies

Dependency	Required By	Impact if Delayed
System access credentials	Week 2	2-week delay per week late
Chairman availability for persona training	Month 1-3	Persona quality degradation
Security requirements sign-off	Month 2	Security rework, timeline risk
UAT resources and participation	Month 5-6	Go-live delay
Production environment access	Month 6	Deployment delay

## 11. Commercial Proposal

### 11.1 Platform Licensing Fee

This section outlines the software licensing cost of Dynamiq's Agentic OS for the client in a private cloud. It includes full access to the Dynamiq platform hosted on the **client's infrastructure**, continuous feature evolution, security patches, and priority support, allocated engineers for support.

Component	Unit	Cost (USD)	Description
Dynamiq Platform Software License	Annually	\$435 967	<ul style="list-style-type: none"><li>- Licensed users: up to 100 developer seats.</li><li>- Monthly requests (agent/workflow executions): up to 10M/month.</li><li>- Agents/workflows deployed: up to 1,000.</li><li>- Vector &amp; graph databases: up to 10.</li><li>- Knowledge bases (RAG indexes): up to 200.</li><li>- Single Sign-On (SSO): custom integration (SAML 2.0 / OIDC).</li><li>- Compliance &amp; privacy: SOC 2, GDPR, HIPAA.</li><li>- Deployment options: cloud, on-premise, or hybrid.</li><li>- Monitoring &amp; analytics: live dashboard (requests, tokens, latency, cost); deep tracing; searchable request logs.</li><li>- Quality &amp; safety: evaluation suite for regression testing; content/PII guardrails; jailbreak/toxicity detection; structured-output enforcement.</li><li>- Model access: multi-model routing via a single API; custom/fine-tuned model hosting.</li><li>- Voice agents support.</li></ul>
24/7 Annual Support	Annually	\$53 437	24x7 premium support with guaranteed SLAs.
<b>Total 1-st Year Investment:</b>		<b>\$489 404</b>	Complete turnkey solution

**Payment Terms:** Annual license billed upfront; services billed at completion/milestones. Prices exclude taxes/duties. **Validity:** 60 days from proposal date.

## 11.2 Implementation Cost

This section details the cost breakdown for the six-month, two-phase implementation and delivery of the Agent-0 platform. It represents the professional services fees for the dedicated Dynamiq team to deploy, configure, integrate, and launch the solution on DGE's infrastructure. The total implementation cost covers all labor from project start through final production sign-off and knowledge transfer.

Phase 1: Foundation & POC (Months 1-3)				
Role	FTE	Days On Project	Day Rate, USD	Total, USD
Project Director, Strategy	1	60	\$1,600	\$96,000
Principal Technical Architect	1	60	\$1,600	\$96,000
Lead AI Engineer	1	60	\$1,600	\$96,000
Senior AI Engineer	2	60	\$1,350	\$162,000
AI Engineer	4	60	\$1,150	\$276,000
Senior Data Engineer	2	60	\$1,150	\$138,000
Data Engineer	2	60	\$800	\$96,000
Backend Developer	3	60	\$800	\$144,000
Frontend Developer	3	60	\$800	\$144,000
DevOps Engineer	1	60	\$1,150	\$69,000
Security Engineer	1	60	\$1,150	\$69,000
QA Engineer	1	60	\$800	\$48,000
Business Analyst	2	60	\$800	\$96,000
Project Manager	1	60	\$1,150	\$69,000
<b>Phase 1 Total:</b>	<b>25</b>			<b>\$1,599,000</b>
Phase 2: Full Build (Months 4-6)				
Role	FTE	Days On Project	Day Rate, USD	Total, USD
Project Director, Strategy	1	60	\$1,600	\$96,000
Principal Technical Architect	1	60	\$1,600	\$96,000
Lead AI Engineer	1	60	\$1,600	\$96,000
Senior AI Engineer	1	60	\$1,350	\$81,000
AI Engineer	3	60	\$1,150	\$207,000
Senior Data Engineer	1	60	\$1,150	\$69,000
Data Engineer	2	60	\$800	\$96,000
Backend Developer	3	60	\$800	\$144,000

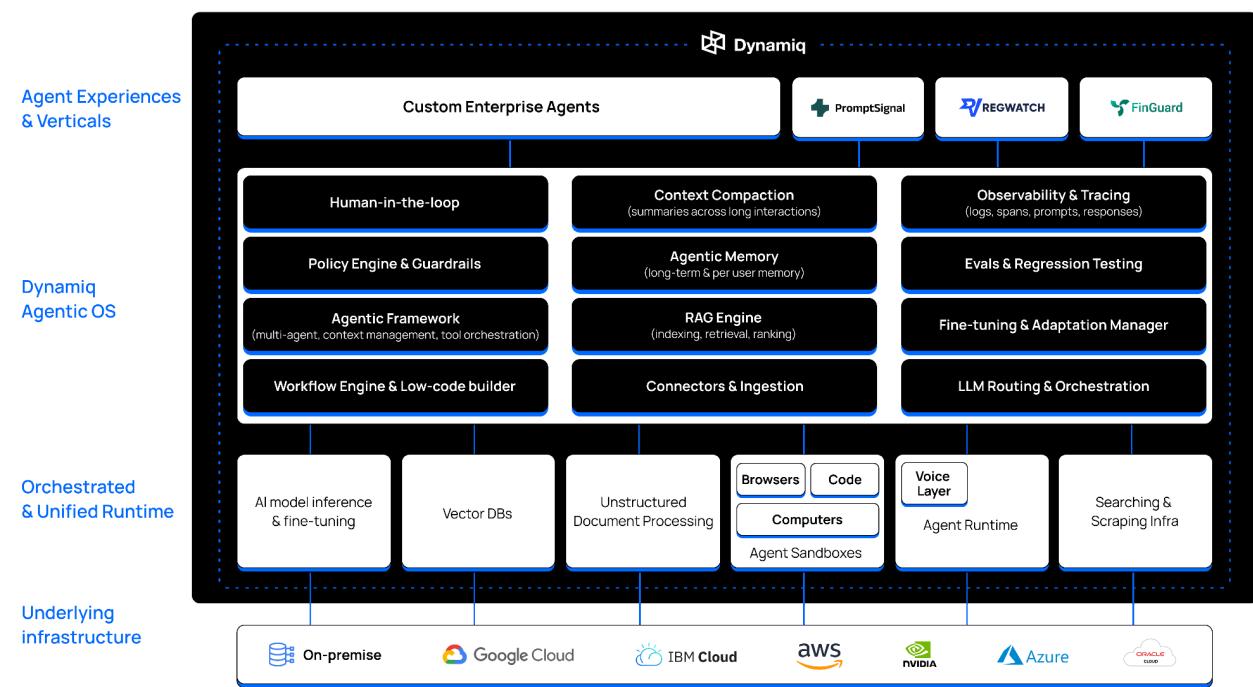
Frontend Developer	2	60	\$800	\$96,000
DevOps Engineer	1	60	\$1,150	\$69,000
Security Engineer	1	60	\$1,150	\$69,000
QA Engineer	2	60	\$800	\$96,000
UX Designer	1	60	\$800	\$48,000
Business Analyst	2	60	\$800	\$96,000
Project Manager	1	60	\$1,150	\$69,000
<b>Phase 2 Total:</b>	<b>23</b>			<b>\$1,428,000</b>
			<b>TOTAL:</b>	<b>\$3,027,000</b>

## Appendix A – Dynamiq Technical Capabilities

Dynamiq provides a sovereign-first agentic AI operating system that enables DGE to deliver Agent-0 as a secure, auditable, cross-government capability deployable in Abu Dhabi government-controlled infrastructure. The platform unifies the core building blocks required for production-grade agents: orchestration, knowledge and RAG, governance guardrails, evaluations/observability, integration, and deployment tooling so that Agent-0 can operate reliably across departments, data sources, and stakeholder workflows.

### 12.1 Sovereign AI Infrastructure & Deployment Flexibility

Dynamiq is **built for sovereign AI deployments**, giving clients full control over their data and models. It can be deployed on-premises or in hybrid environments, including air-gapped data centers, and is aligned with strict security standards such as SOC 2 and GDPR. All AI processing runs inside the customer's environment so sensitive data does not leave the sovereign boundary.



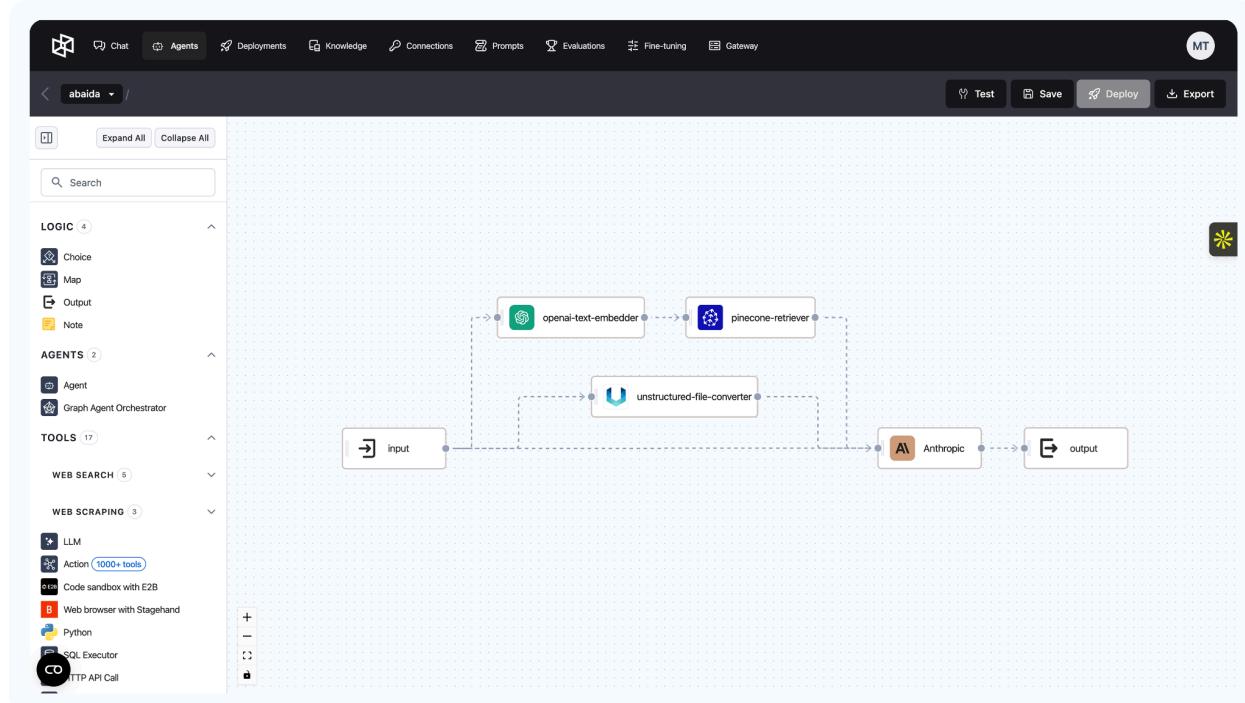
**Figure 12.1 Sovereign agentic OS stack:** deploy Dynamiq on-prem, hybrid, or air-gapped with all AI workloads running inside the customer boundary.

The platform is **cloud-agnostic** and containerized, running on Kubernetes and Docker across on-prem servers or major clouds (AWS, Azure, Google, IBM, Oracle). This supports existing GPU fleets (e.g., NVIDIA L40, A100, H100) and hybrid scaling when needed. The result is a dedicated AI stack that meets data-residency requirements and scales to enterprise workloads.

Dynamiq supports a broad range of open-source LLMs and supporting components. It can deploy a variety of AI models, selecting the right model size per workload. Larger models are served efficiently using distributed inference techniques (e.g., sharding and batching). The platform includes a managed fine-tuning pipeline, data preparation through training and deployment, using parameter-efficient methods such as LoRA to reduce resource needs. This allows organizations to run and adapt leading open models (e.g., Llama, DeepSeek, Qwen) within their own environment. Dynamiq also supports open-source vector databases (e.g., Qdrant, Weaviate) to enable fully sovereign, air-gapped retrieval systems.

## 12.2 Low-Code Workflow & Agent Builder

Dynamiq provides a **low-code interface** for designing AI agent workflows. With a drag-and-drop canvas, teams assemble pipelines; data ingestion, RAG, LLM calls, tool integrations, and outputs, with minimal code. Workflows are built from nodes representing actions or agents, connected to define the process. Common building blocks (database queries, API calls, document loaders) are included so teams can ship GenAI applications faster. For example, an agent can ingest a PDF, run an LLM analysis, and route results into a business system without boilerplate development.



**Figure 12.2 Dynamiq low-code workflow builder:** Drag-and-drop canvas to compose agent workflows (inputs, RAG/embedding, tools, LLM calls, orchestration, and outputs) with built-in testing and versioning.

The platform is also extensible for engineers. Users can insert custom Python blocks or bring custom components into the same workflow. Dynamiq includes hundreds of connectors to enterprise systems (e.g., Jira, SharePoint, Google Drive, email, databases) and supports

**1,000+ pre-built integrations** for ingestion, tool calling, and API-driven workflows. Teams can prototype in a sandbox, then promote workflows to production with controlled rollout. The studio includes testing and debugging so users can step through execution, refine prompts, and tune parameters before deployment.

## 12.3 Advanced RAG and Knowledge Management

Dynamiq provides a Knowledge Hub to ingest documents and data from enterprise sources into an AI-ready knowledge base. Users connect repositories and systems (Google Drive, SharePoint, Box, Confluence, websites, databases) via built-in connectors. Documents are automatically chunked and embedded using the selected embedding model. Dynamiq supports multiple vector stores (pgVector, Pinecone, Weaviate, Qdrant) so organizations can choose the database that fits their infrastructure and scale. The RAG pipeline is managed end-to-end: once ingested, content becomes securely queryable by agents.

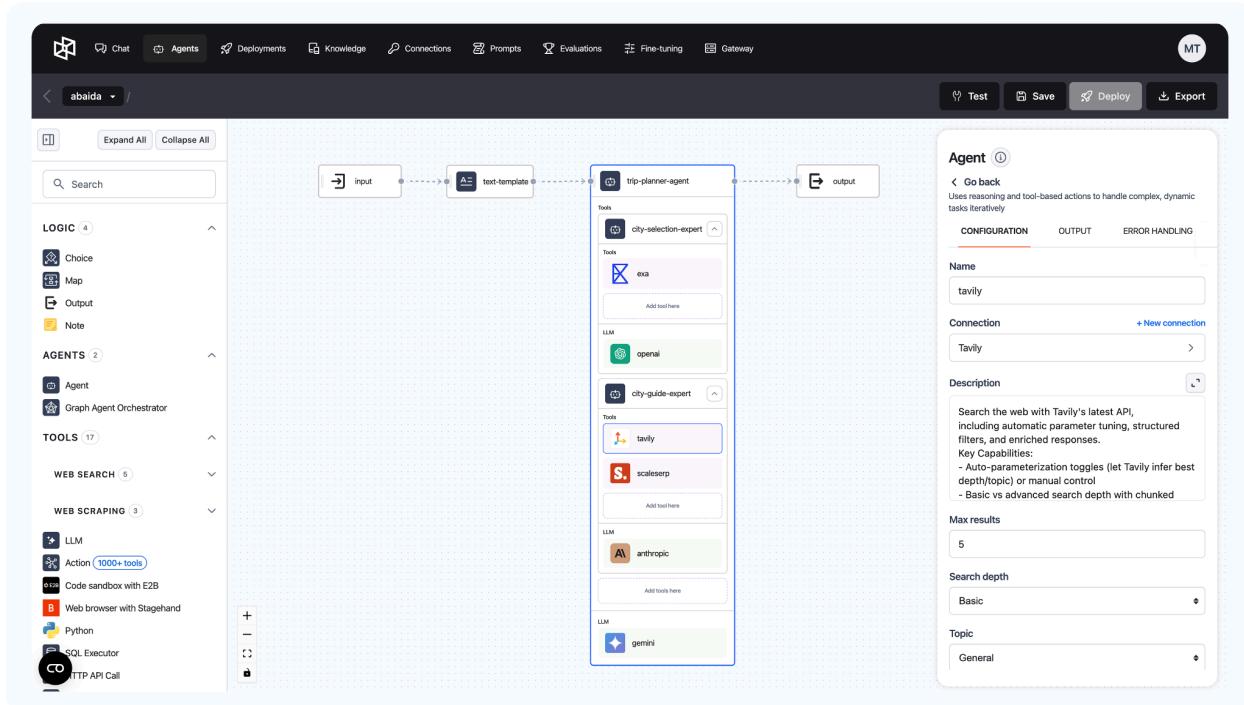
The screenshot shows the Dynamiq Knowledge Base interface. At the top, there's a navigation bar with links for Chat, Agents, Deployments, Knowledge, Connections, Prompts, Evaluations, Fine-tuning, and Gateway. A user profile icon 'MT' is in the top right. Below the navigation is a breadcrumb path: 'abaida' > 'Knowledge bases' > 'healthcare-assistant'. The main area displays a card for the 'healthcare-assistant' knowledge base, showing its Ingestion Endpoint (fabbb6cd-de25-4bfe-ba6d-5845296c4786...), Runtime (1281), Ingestion Workflow (healthcare-assistant v1), Last edited by (MT - Maria-Elena Tzanev at Dec 18, 2025 3:44 PM), and Deployed by (MT - Maria-Elena Tzanev at Dec 18, 2025 3:44 PM). Below this, there's a tab bar with FILES, WORKFLOW, INTEGRATIONS (which is active and highlighted in red), INGESTION, and RETRIEVAL. The INTEGRATIONS section lists various enterprise content sources with 'Connect' buttons: Google Drive, Notion, Dropbox, Microsoft OneDrive, Microsoft SharePoint, Box, and Website. Each item has a small icon to its left.

**Figure 12.3 Dynamiq Knowledge Base integrations:** connect common enterprise content sources to ingest and maintain up-to-date knowledge for secure RAG-powered agents.

Agents use RAG nodes in workflows to retrieve relevant context during conversations. Data access policies are enforced at retrieval time, so agents only access permitted content and cross-department leakage is prevented. Teams can tune retrieval by configuring the vector backend, embedding model, chunking strategy, and similarity thresholds. This enables grounded, current answers based on enterprise sources (policies, PDFs, emails, and more) while keeping knowledge and queries inside the sovereign environment and governed by the same compliance controls as the rest of the platform.

## 12.4 Complex Multi-Agent Orchestration and Tool Automation

Dynamiq's orchestration engine manages complex AI workflows by coordinating multiple specialized agents<sup>[101]</sup>. Instead of a single-step interaction, a task is broken into sub-tasks that run sequentially or in parallel. For example, one agent might gather data while another synthesizes it into a report; Dynamiq seamlessly passes context and interim results between them. The platform provides built-in Linear, Adaptive, and Graph workflow patterns that let teams set up fixed sequences, decision-based branches, or parallel flows as needed. This flexibility means even multi-step processes, like cross-agency approvals or detailed analyses, can be automated end-to-end without custom coding.



**Figure 12.4 Multi-agent workflow orchestration:** visual canvas to compose role-based agents (planner and specialist agents), connect tools and LLMs, and route outputs through a governed multi-step flow from input to final response.

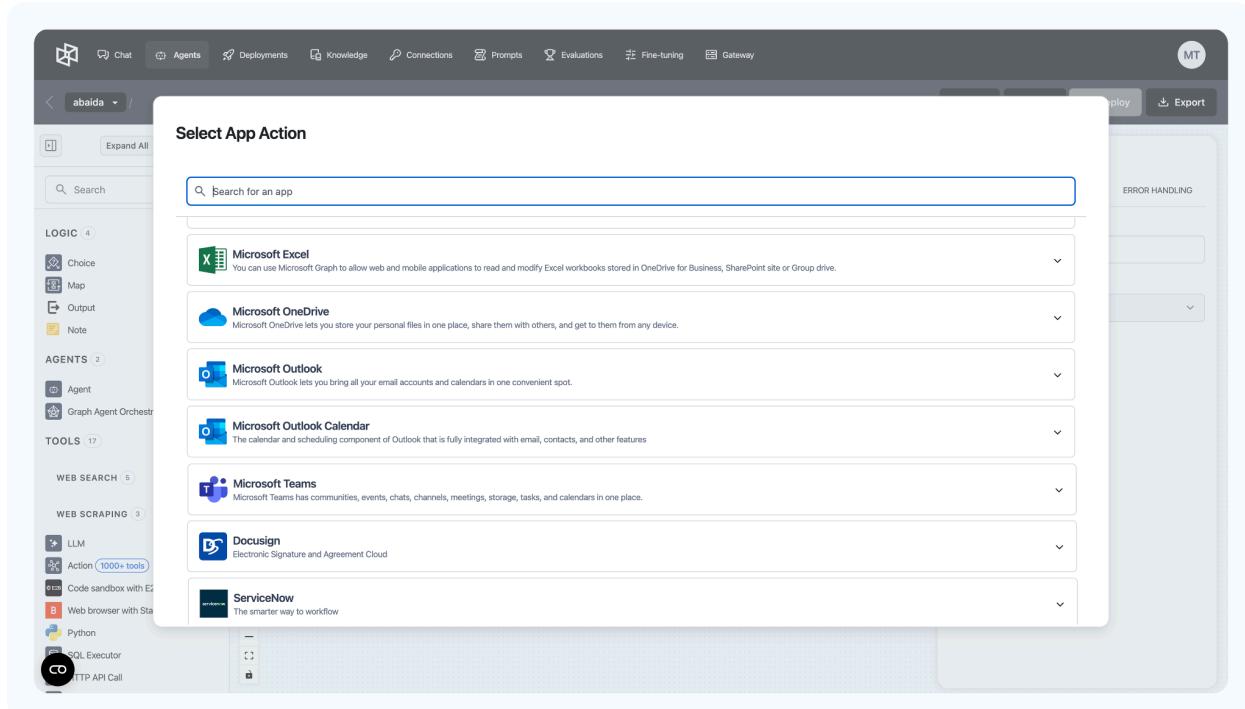
Multiple agents work like a specialized “AI team”, each contributing its expertise<sup>[101]</sup>. For example, one agent could focus on fact-finding, another on drafting a response, and a third on verification;

the engine automatically hands off context and merges their outputs. Built-in checkpoints allow human review at any stage, keeping critical workflows governed. This multi-agent design boosts accuracy and speed, and since every action is logged, it provides a complete audit trail. In practice, that means higher quality results: each agent handles the work it's best at, and Dynamiq's coordination ensures the pieces fit together reliably.

For DGE's Agent-0, this orchestration approach makes cross-agency tasks seamless and auditable, aligning with Dynamiq's goal of letting Agent-0 operate reliably across departments and data sources [OBJ]. For example, Agent-0 could generate a government-wide report by having one agent gather policies from different ministries while another analyzes budget data; Dynamiq would automatically coordinate these subtasks. A final agent would then compile the findings into a unified answer. Because each step is managed and logged, the result is faster and more reliable, and officials can audit any part of the process. Dynamiq essentially turns Agent-0 into a dependable "AI team" that accelerates insights across government while ensuring compliance.

## 12.5 Native, Governed Tool Calling with Full Auditability

Beyond orchestration, the platform natively integrates external tools and services into agent workflows. Agents can invoke web services, databases, enterprise APIs, and other external systems through a unified API-as-a-tool interface, treating any **API as a tool** that can be registered, permissioned, and called during execution [OBJ]. Crucially, all such tool usage is governed and auditable: whenever an agent calls a tool, the platform handles the invocation and returns the result to the agent, while **logging the action for complete auditability**. This standardized tool-calling mechanism expands an agent's capabilities from answering questions to performing transactions or updates in other software, all under enterprise control and oversight [OBJ].

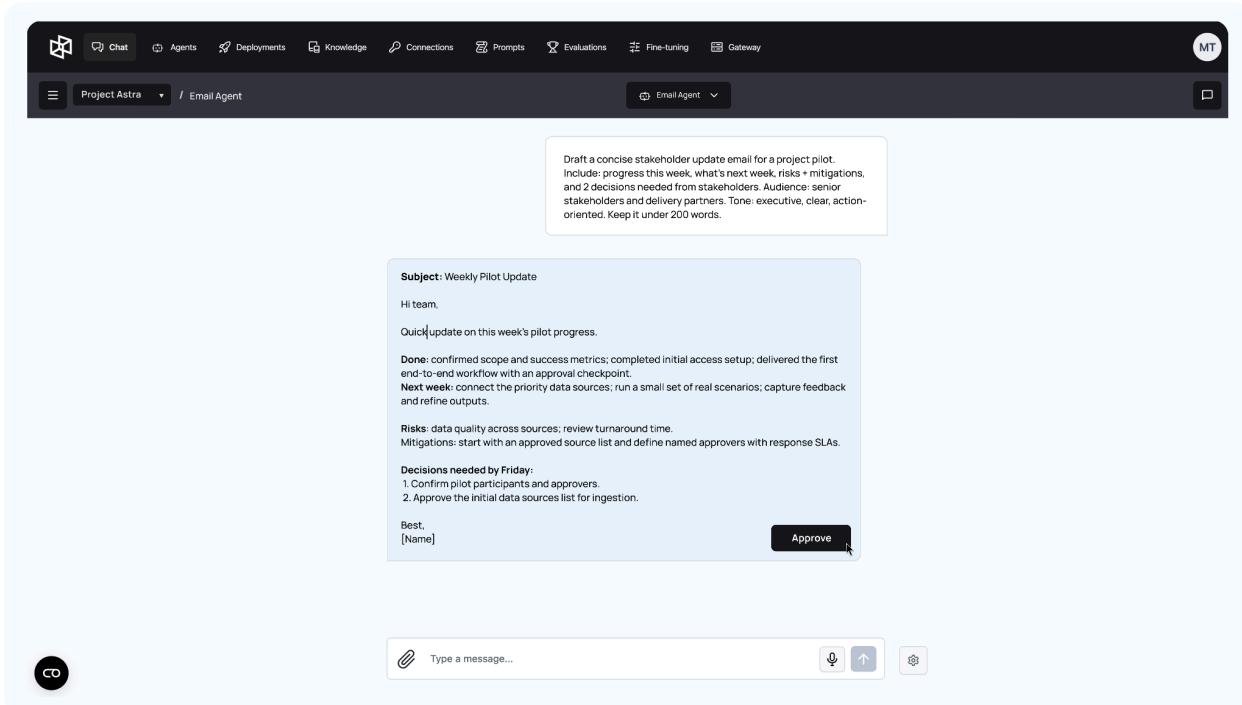


**Figure 12.5 Native, governed tool calling:** Dynamiq lets agents securely invoke approved enterprise apps and actions through a controlled tool layer with permissioning and auditability.

Each agent operates under granular, scoped permissions, and every tool invocation carries a unique identity, an approach that provides clear **audit trails for all agent actions**. Every tool call is recorded as a structured log entry (with correlation IDs, timestamps, inputs/outputs), creating an immutable record that ensures compliance and facilitates forensic analysis of agent-driven operations. By combining delegated authentication with fine-grained logging, the system guarantees that tool usage remains secure, transparent, and fully accountable. In short, agents can take action through external tools in real time, but nothing happens “in the dark” – every step is permissioned and traceable end-to-end.

## 12.6 Human-in-the-Loop

For high-stakes or sensitive tasks, the platform provides built-in **human-in-the-loop** controls to keep a human decision-maker involved at critical junctures. Workflow designers can insert **approval checkpoints** into an agent’s plan so that certain actions or decisions pause and await human review before proceeding. Through the platform’s interface, designated supervisors or subject-matter experts can inspect an agent’s reasoning and proposed action at that point and then choose to approve, reject, or modify the agent’s step as needed. These approval steps can be configured without any coding in the visual builder, making it straightforward to embed human oversight into automated agent processes whenever required.

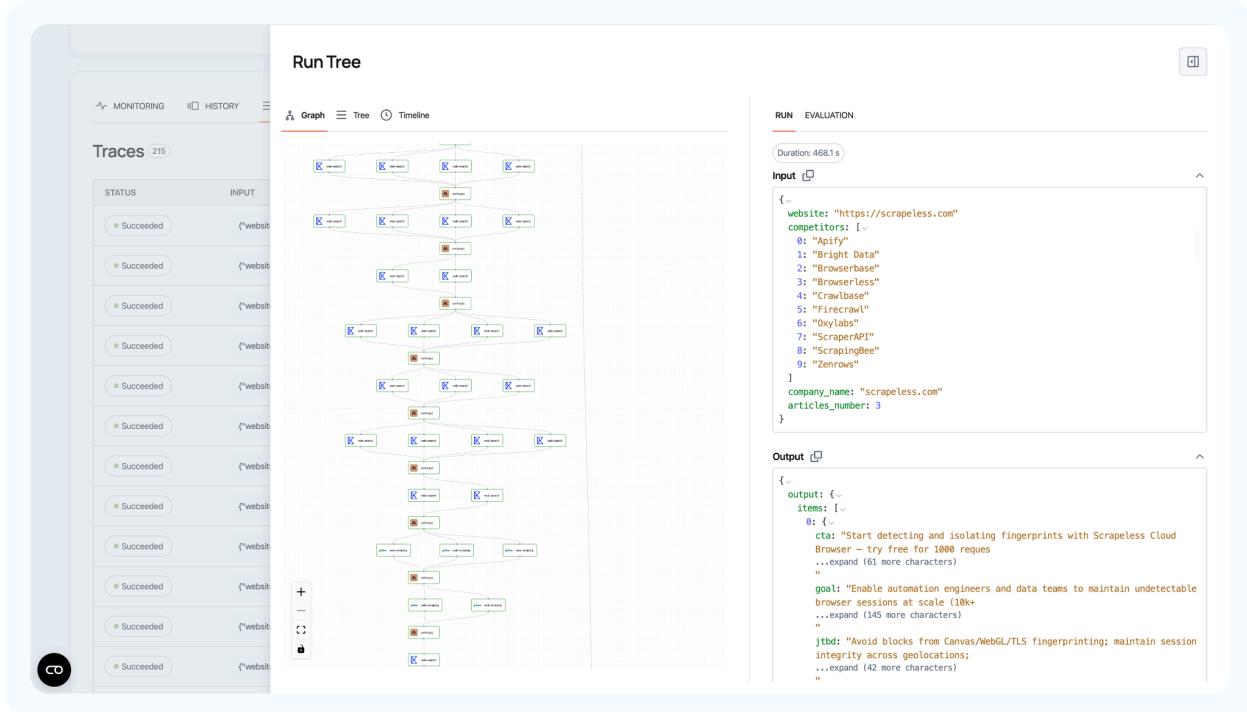


**Figure 12.6 Human-in-the-loop approval for high-impact outputs:** Dynamiq routes agent-generated drafts (e.g., stakeholder emails) through a review gate where a user can edit and explicitly approve before anything is finalized or sent.

This capability ensures that **sensitive operations always have human sign-off**. Critical decisions (for example, sending out an official communication or making a significant financial transaction) can be double-checked by a person, preserving accountability and alignment with organizational policies. The agent will only continue its execution after receiving the necessary human approval, or it may adjust its course based on human feedback. As a result, the organization maintains final authority over important outcomes: even though AI agents handle the heavy lifting of analysis and execution, **ultimate control remains with human experts**, preventing the AI from overstepping bounds. By blending AI automation with human judgment at key points, the platform enables scaling of processes and decision-making while ensuring that oversight and responsibility are never lost. This human-in-the-loop approach builds trust in the system's outputs and guards against unintended consequences in sensitive scenarios.

## 12.7 Comprehensive Agent Evaluation & Observability

Dynamiq includes observability and evaluation capabilities for operating agents in production. Every agent run can be logged and traced, including token-level histories with timestamps and metadata. This supports debugging, root-cause analysis, and explainability. Dynamiq tracks usage metrics (token counts, latency, model costs) and provides spend controls to cap usage and prevent runaway costs. Telemetry can be exported via OpenTelemetry to integrate with existing monitoring stacks.

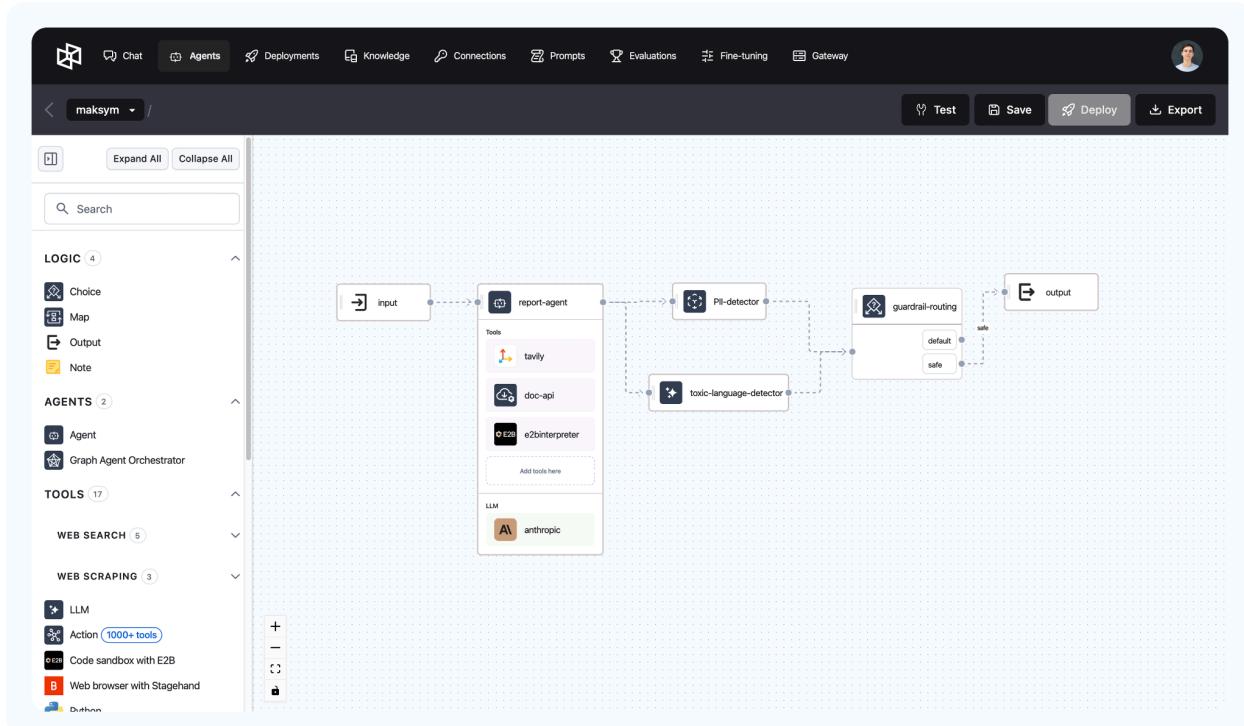


**Figure 12.7 Comprehensive agent evaluation and traceability:** Dynamiq provides end-to-end run traces with step-by-step graphs, inputs/outputs, and evaluation views to validate agent behavior, debug workflows, and support audit-ready governance.

Dynamiq also includes evaluation tooling to test and validate agents before deployment and after changes. Teams can define evaluation criteria (accuracy vs. ground truth, format checks, policy compliance) and run offline test suites to detect regressions. Evaluation results are surfaced in the UI to support release readiness. User feedback can be captured, incorporated into metrics, and used to guide updates. A visual trace viewer provides end-to-end inspection of prompts, tool calls, and outputs, supporting audits and continuous improvement.

## 12.8 Execution Guardrails and Enterprise Policy Enforcement

The platform is designed with rigorous guardrails and policy enforcement mechanisms to ensure AI agents operate safely within enterprise and regulatory boundaries. Robust security and governance controls are applied across the board. For instance, role-based access control (RBAC) is enforced with fine-grained permissions on who can build, deploy, or execute agents, and **immutable audit logs** capture every agent action (model queries, tool calls, data access, etc.) along with the identity and timestamp<sup>[OBJ]</sup>. All agent activities are thus traceable and attributable, supporting comprehensive audits. Administrators can also define explicit policies that constrain agent behavior – for example, requiring human approval for certain categories of requests or outright blocking disallowed operations – and the platform will automatically enforce these rules during agent execution<sup>[OBJ]</sup>. This ensures that agents stay within the **bounds of organizational policies**, no matter how autonomous their reasoning might be.



**Figure 12.8 Guardrails and policy-based routing.** Dynamiq applies automated checks and routes requests through a guardrail decision layer before returning outputs, enabling safe, compliant agent execution.

Moreover, Dynamiq includes a suite of AI-specific guardrails to maintain compliance and safety in generated outputs. The system can detect and redact personally identifiable information (PII) in agent outputs, apply prompt-injection filters to sanitize inputs, and perform content moderation to filter out toxic or policy-violating language, so that the AI's responses meet enterprise compliance requirements (e.g. adhering to SOC 2 and HIPAA standards) [obj]. Developers can even enforce structured output formats (for instance, requiring an agent's answer to conform to a JSON schema) to prevent unpredictable free-form results and to simplify downstream integrations [obj]. Additional operational guardrails like execution timeouts, step limits, and usage quotas can be configured to prevent runaway processes or resource misuse, further bounding the agent's autonomy within safe limits [obj]. All these controls are configurable to align with the organization's internal policies and industry regulations, and they work in tandem to **keep the AI's actions transparent, predictable, and compliant**.

These guardrails and enforcement features are critical for deploying AI in sensitive or regulated environments. They give the organization confidence that even as agents plan and act independently, they **remain compliant, auditable, and under control** at all times [obj]. The platform effectively creates a governed sandbox for agent operations: the AI can be innovative and proactive in achieving its goals, but it cannot violate the rules, thresholds, or ethical guidelines defined by the enterprise. By enforcing strict policies and safety checks at every step of execution, Dynamiq enables scaled-up agentic automation **without sacrificing oversight or governance**. This means an enterprise can embrace powerful AI-driven workflows and decision-making, while ensuring security, accountability, and trust are never compromised.

## 12.9 Open-Source LLM Deployment and Fine-Tuning

Dynamiq enables organizations to deploy and operate open-source LLMs and customize them for domain needs. The platform supports model selection, deployment, fine-tuning, and production monitoring within one interface. Models can run on-prem for sensitive workloads or in approved cloud environments, and can be integrated into workflows alongside managed services (e.g., Amazon Bedrock) when appropriate.

The screenshot shows the 'Add new deployment' interface in Dynamiq. On the left, there's a sidebar with 'Advanced configuration' expanded, showing 'Replica Autoscaling' settings (Min: 1, Max: 10), 'Max Model Length' set to 4096, and 'Max Number of Batched Tokens' set to 8192. Below that is the 'Quantization' section with 'fp8' selected. The main area has fields for 'Name' (Enter name), 'Description' (Enter Description), 'Model' (Select model), 'Engine' (vLLM selected), 'Runtime' (vLLM (0.9.1)), and 'Resource profile' (aws p4de.24xlarge: 8x NVIDIA A100, Intel, 96 vCPUs, 115). At the bottom are 'Back' and 'Create' buttons.

**Figure 12.9 Open-source LLM deployment with autoscaling and performance controls:** Dynamiq lets teams deploy and configure LLM endpoints with replica autoscaling, context and batch settings, and quantization options for efficient, production-grade inference.

Dynamiq provides a fine-tuning engine (PyTorch-based) that supports LoRA adapters or full fine-tunes with minimal setup. The platform orchestrates preprocessing, training execution (on local GPU clusters or cloud), model versioning, and deployment to inference servers. Fine-tuned models can be A/B tested, rolled back, and monitored. Vector databases and embedding models can be deployed as part of the same stack, enabling a private, enterprise-grade “ChatGPT-like” capability on proprietary data with full model and data control.

**Add new fine-tuned model**

Name: deepseek-financial-data

Model: deepseek-ai/DeepSeek-R1-0528

Resource profile: aws p4de.24xlarge: 8x NVIDIA A100, Intel, 96 vCPUs, 115: ^

Description: Enter description

Hyperparameters:

- No. of epochs: 10
- Batch size: 4
- Learning rate: 0.0001
- LoRA rank: 16

Cancel Next

**Add new fine-tuned model**

Training data [Sample JSONL](#)

Upload a JSONL file with essential data for fine-tuning, including input prompts and desired outputs. Refer to the sample file for guidance and ensure adherence to the specified structure.

Click to upload or drag and drop JSONL file

Cancel Create

**Figure 12.10 No-code fine-tuning workflow for domain adaptation.** Dynamiq enables teams to create and version fine-tuned models by selecting a base LLM, uploading training data (JSONL), and configuring key hyperparameters for rapid, on-infrastructure customization.

## 12.10 API-First Integration Architecture

Every capability within the platform is accessible through well-documented APIs that support multiple integration patterns. REST APIs provide standard request-response interactions for system integration. WebSocket connections enable real-time bidirectional communication essential for interactive agent conversations. Server-sent events deliver streaming responses for long-running analyses, allowing stakeholders to monitor progress in real-time.

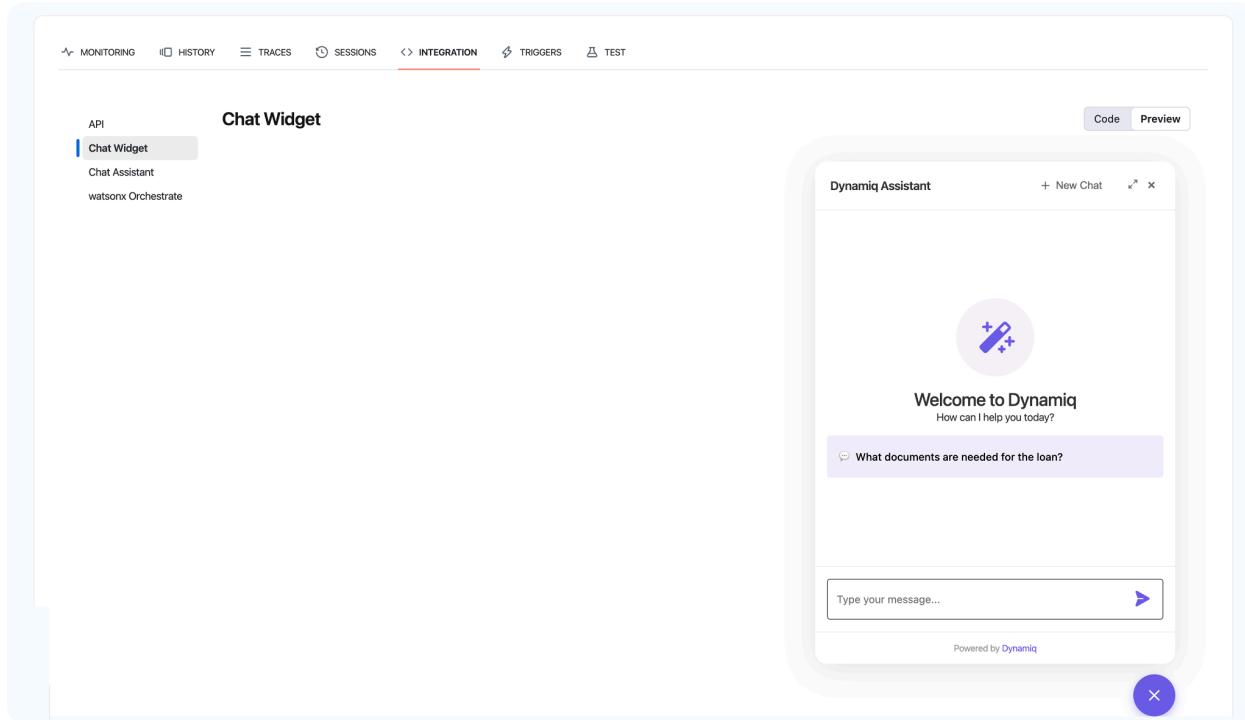
The screenshot shows the Dynamiq platform's API integration interface. At the top, there are navigation tabs: MONITORING, HISTORY, TRACES, SESSIONS, INTEGRATION (which is highlighted), and TEST. Below the tabs, there are three main sections: Chat Widget, watsonx Orchestrate, and Chat Assistant. The Chat Assistant section is expanded, showing a code example for "Server-Sent Events (SSE) Streaming". The code is written in Python and demonstrates how to send a payload to a specific endpoint using SSE streaming. The code includes imports for os, requests, and json, defines an endpoint URL, retrieves an access token from environment variables, sets headers for Content-Type and Authorization, and creates a payload with an input schema and a stream flag. The Python tab is selected, and there is also a TypeScript tab.

```
1 import os
2 import requests
3 import json
4
5 endpoint = "https://f10664d4-06f0-4bcd-a3f8-259650b3de57.apps.us-east-1.aws.getdynamiq.ai"
6 token = os.getenv("DYNAMIQ_ACCESS_KEY") # Generate Access Key in the UI settings
7
8 streaming_event = "data" # Event name configured in the UI
9
10 headers = {
11     'Content-Type': 'application/json',
12     'Authorization': f'Bearer {token}',
13 }
14
15 # Payload: Modify input schema as per the input node schema defined in the UI
16 payload = {
17     "input": {
18         "input": "Explain basic Machine Learning concepts."
19     },
20     "stream": True,
21 }
22
```

**Figure 12.11 API-first integration and streaming.** Dynamiq exposes secure REST endpoints with Server-Sent Events (SSE) streaming and ready-to-use Python/TypeScript examples, enabling fast integration of agents into existing applications and portals.

## 12.11 Pre-Built UI Components and LLM Gateway

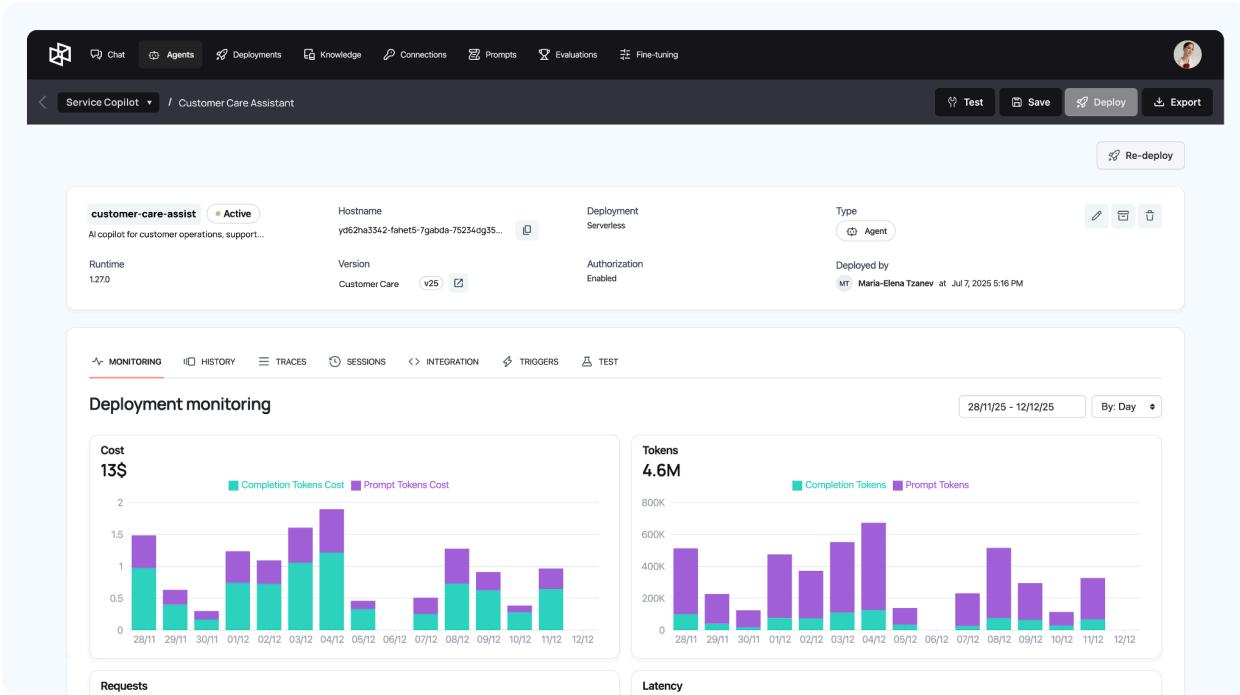
Dynamiq includes a comprehensive library of pre-built UI components and APIs that accelerate delivery of agent experiences without dedicated frontend development. These components are delivered as drop-in widgets and can be embedded into any web application, SharePoint site, or internal portal using lightweight JavaScript integration. All components are fully white-labeled and configurable to match organizational branding and UX standards.



**Figure 12.12 Built-in, white-label chat widget for rapid rollout:** Dynamiq provides an embeddable chat UI to deploy agents to end-users quickly, with configurable branding and interaction patterns.

The core component is the intelligent chat widget, which provides a complete conversational interface for interacting with agents. It supports multi-modal inputs (text, voice, file uploads) and delivers rich outputs, including tables, charts, and interactive visualizations. The widget maintains searchable conversation history, can suggest follow-up questions to guide users toward productive interactions, and supports Arabic and English. It also includes configurable behavior controls to align interactions with organizational policies and preferred response patterns.

Beyond chat, Dynamiq provides purpose-built components for common workflows. The document analysis widget enables drag-and-drop document processing with real-time extraction and summarization. The dashboard widget consolidates metrics and insights from multiple agents into executive-ready visualizations. The low-code workflow and AI agent builder widget provides a visual interface for non-technical users to create and update agent workflows without writing code. Components connect securely via WebSocket or Server-Sent Events (SSE), enabling real-time updates without page refreshes.

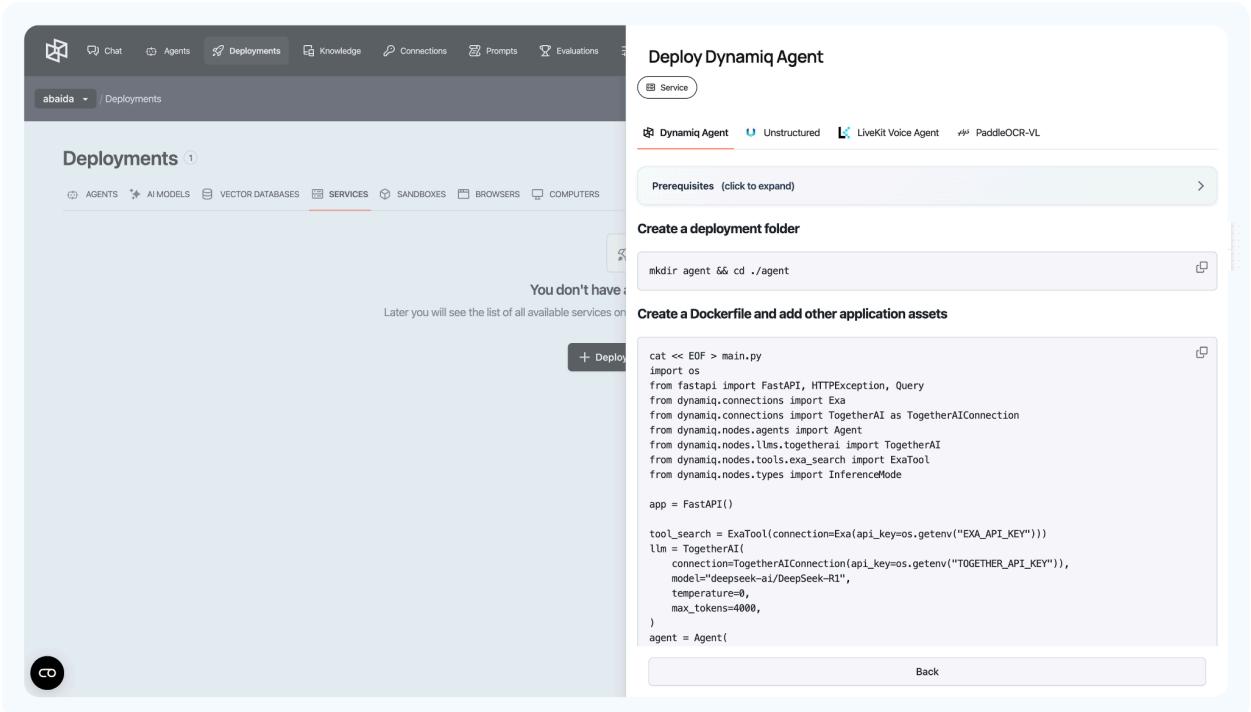


**Figure 12.13 Built-in observability and monitoring dashboards.** Dynamiq provides deployment-level monitoring for agents and workflows, including token usage, cost, latency, and run history, enabling operational oversight and governance in production.

Dynamiq also provides a unified LLM Gateway that exposes a single REST/WebSocket interface for hosted and external models. The gateway routes requests to the selected model or agent and handles authentication, rate limiting, and load balancing. It supports streaming responses over WebSocket/SSE for interactive user experiences and consistent integration patterns across applications. A unified chat console allows developers to test and compare deployed models and agents, accelerating prompt iteration, validation, and controlled rollout.

## 12.12 Open Architecture and Extensibility

Dynamiq is built on an open-source agent framework (Apache 2.0) and supports both low-code and code-first workflows. Visual workflows can be represented as code for version control, automation, and extensibility. The platform is infrastructure-agnostic and runs across hyperscalers or bare-metal, supporting multi-cloud strategies and sovereign deployments. Dynamiq provides deployment tooling (e.g., Terraform, Helm, CI/CD integrations) so the platform fits standard enterprise DevOps practices.



**Figure 12.14 Open, extensible deployment model:** Dynamiq packages agents as standard services and deploys them into customer infrastructure, enabling “bring-your-own” code, tools, and models without vendor lock-in.

Dynamiq also aligns with emerging standards and ecosystem integrations. It contributes to open patterns for agent-tool connectivity (MCP) and integrates with major AI and data ecosystems. Customers can export models, indexes, and logs in standard formats, avoiding vendor lock-in and enabling long-term portability.

# Appendix B – Dynamiq Case Studies & Tutorials

## 1. Case Studies:

- a. [Smart Document Search for a B2B Accounting Platform Built in One Week with Dynamiq](#)
- b. [Building Better Cancer Treatment Plans with Dynamiq's Multi-Agent AI](#)
- c. [Automating Customer Support at Scale: How a Neo Bank Saves \\$1.5M a Year with AI](#)

## 2. Tutorials:

- a. [Build an Intelligent Agent System for Market Analysis with DeepSeek](#)
- b. [Building a Search GPT with Dynamiq](#)
- c. [Build an Autonomous Data Analyst using Dynamiq, E2B and Together.ai](#)
- d. [Agent Orchestration Patterns in Multi-Agent Systems: Linear and Adaptive Approaches with Dynamiq](#)
- e. [Building a Conversational AI Agent with Dynamiq](#)
- f. [Automating customer support in banking using agentic RAG](#)
- g. [Building an Intelligent Linear App Assistant](#)
- h. [Build an Automated Social Media Management Agent with Dynamiq](#)

## Appendix C – Dynamiq Infrastructure Components Info

#	Component Name	Component Description	Tech Stack
1	Dynamiq OSS	Open-source orchestration framework for agentic AI and LLM applications	Python
2	UI	Front-end service for the agentic AI platform	React SPA, Typescript
3	Core API gateway	Main API gateway for agentic AI platform for mission critical and latency-sensitive endpoints	Golang
4	AI-native API gateway	Secondary API gateway for AI-native and Python-specific workloads (prompt testing, LLM chat, testing of AI agents, agent evaluations)	Python, FastAPI
5	Agentic runtime	Microservice that allows deploying and running in a serverless mode	K8S + Karpenter + Python, FastAPI
6	Custom-code service runtime	An environment that allows engineers deploying any Python/Docker based agents and services without the no-code builder	Golang + Native Kubernetes wrapper
7	LLM fine-tuning service	An environment that allows LoRa-based fine-tuning of open-source LLMs	Python, transformers
8	Multi-LoRa inference service	Microservice that allows running LLM models with LoRa layers	vLLM, LoRaX, Python, Kubernetes
9	LLM Embedding Serving	A service that enables open-source embedding models inference	Python, transformers, ONNX, FatAPI
10	Pre-built Agentic UI components	UI components that allow users to embed AI agents into any web-application without a code	React, JavaScript
11	Infrastructure-as-a-Code and CI/CD	All Terraform code, Helm charts, ArgoCD Pipelines, Github Actions	Terraform, Helm, Python, Bash
12	Platform documentation	Entire white-labeled documentation of Dynamiq's infrastructure	Markdown